

Sreevarshini Srinivasan

Data Engineer

College Park, MD | varshini.sree03@gmail.com | LinkedIn | GitHub | Portfolio

Data Engineer experienced building Python-based ETL/ELT pipelines and cloud platforms (AWS, Azure, Databricks) for regulated, finance-adjacent enterprise systems. Skilled in Unix/Python scripting, CI/CD, and production monitoring, delivering reliable, scalable solutions across cross-functional teams.

EXPERIENCE

Loyaltytics AI

Data Engineer

Apr 2023 – Aug 2024

Bengaluru, India

- Built and scaled Python/PySpark ETL pipelines (Databricks Delta Live Tables) processing **15GB+ daily** across **10+ sources and 100+ tables** for regulated retail and financial-reporting systems.
- Developed an internal tax-calculation platform processing **6 years of historical financial data** across **7 countries** with **100% accuracy**, supporting regulated financial reporting requirements.
- Built automated data validation and testing frameworks in a Linux environment, reducing storage costs by **70%** and improving pipeline reliability across production systems.

DriverAI, LLC

Cloud Infrastructure Engineer / Intern

Jun 2025 – Aug 2025

Remote

- Provisioned and automated AWS (EC2, RDS) and Azure infrastructure using Terraform, reducing manual deployment overhead by **60%** and improving reproducibility across environments.
- Built FastAPI/PostgreSQL backend services, supported CI/CD pipelines with Git-based version control, and implemented CloudWatch/Grafana monitoring to track performance and resolve incidents.

Loyaltytics AI

Data Engineering Intern

Jan 2023 – Mar 2023

Bengaluru, India

- Designed Azure Data Factory pipelines orchestrating ingestion from Oracle, MySQL, XML, and CSV sources, authoring SQL/Python transformation and validation logic.
- Automated data validation and testing processes, improving pipeline efficiency by **40%** while supporting Unity Catalog governance migration for a major retail chain.

TECHNICAL SKILLS

Python, SQL, AWS (EC2, S3, RDS, CloudWatch), ETL/ELT Pipelines, PySpark, Databricks, Delta Lake, Unix/Linux, Bash, Git, CI/CD, DevOps, Terraform, Azure, Data Structures & Algorithms, Data Modeling, Data Validation & Testing, Monitoring (Grafana, CloudWatch), Apache NiFi, Apache Airflow, PostgreSQL, Oracle, MySQL, Unity Catalog

PROJECTS

SchemaDrift — Real-Time Schema Drift Detection & Lineage System

Python, Kafka, FastAPI

- Built a metadata-driven platform using graph data structures (Neo4j) to detect schema drift in real time across PostgreSQL, Kafka, S3, and Delta Lake pipelines, enabling automated blast-radius analysis.
- Designed automated validation rules and CI/CD-style blocking logic to catch unsafe schema changes and data integrity issues before downstream production impact.

Feature Store ML Pipeline & Drift Detection System

Python, PySpark, Delta Lake, Airflow

- Built a Python/PySpark ETL pipeline writing versioned feature tables to Delta Lake, orchestrated via Apache Airflow with automated monitoring and alerting for daily pipeline runs.
- Implemented automated drift detection and a quarantine/retraining workflow, improving model reliability through data-informed, continuous-improvement practices.

EDUCATION

University of Maryland, College Park

Aug 2024 – May 2026

Master of Science, Data Science

Relevant Coursework: Algorithms, Probability Theory, Statistical Pattern Recognition, Deep Learning

CERTIFICATIONS

Databricks Certified Data Engineer Associate, Microsoft Certified: Fabric Data Engineer Associate, Google Cloud: Architecting with Compute Engine Specialization